

THE CATHOLIC UNIVERSITY OF EASTERN AFRICA

Faculty of Science
Department of Natural Sciences

DISSERTATION

Measuring the External Field: A Cortical-Proxy Content Observable and the Mean-Field Bound on Media-Driven Opinion Change

*An instrument, a 400-item corpus,
and a source-confounded separation*

Student Name: Brian Mwai
Admission Number: 1050555
Programme: Bachelor of Science in Physics
Supervisor: Dr. Songa Mutambi
Date: September 2026

A dissertation submitted in partial fulfillment of the requirements
for the degree of Bachelor of Science in Physics.

Declaration

This dissertation is my own original work and has not been presented for the award of a degree at this or any other university. The instrument, the corpus, the analysis scripts and the text are mine. Where the work of others is used it is cited, and the released model checkpoint, the public datasets and the software libraries the work builds on are named in the chapters that use them.

No part of the measurement or the analysis was contributed by any person other than the author. The supervision named below was academic guidance and approval, not authorship.

Brian Mwai

Admission Number 1050555

Signature

Date

This dissertation has been submitted for examination with my approval as the university supervisor.

Dr. Songa Mutambi

Department of Natural Sciences,
Faculty of Science
The Catholic University of Eastern
Africa

Signature

Date

Abstract

Media is routinely modelled in statistical physics as an external field acting on a population of coupled opinion states, and that field is almost never measured. This dissertation builds an instrument that measures a candidate observable from content itself, applies it to a corpus, and reports what the measurement can and cannot support.

The observable is a cortical proxy. Text is spoken, transcribed, embedded, and passed to TRIBE v2, a released encoder that predicts cortical response; two networks are averaged from that prediction, one associated with affective salience and one with deliberative control, and the index is their signed difference. The ratio form specified in the proposal is undefined for 17.25% of items and was replaced by that difference. Nothing here is a measurement of a person: no participant was scanned, and the observable is never identified with any subcortical structure.

A corpus of 400 items across four categories, matched on length, was scanned on a GPU. The categories separate at $\eta^2 = 0.1068$, $F = 15.779$, $p = 1.035 \times 10^{-9}$, against a design powered to detect $\eta^2 = 0.0268$, and the index discriminates the pre-scan manipulative label at $AUC = 0.6274$, against 0.9758 for a TF-IDF baseline on the same rows, so the index is reported as a field observable and not as a competitive classifier. The separation is carried almost entirely by fear-activating content, $d = 0.9389$ against neutral, while high-outrage content does not separate at all, $d = 0.0300$, though it is the category the proposal expected to separate most; splitting the index into the two activations it differences shows high outrage raising both by almost the same amount, $d = 0.507$ and $d = 0.491$, so its flat difference is a symmetric response rather than an absent one. Each category is drawn from a single source dataset, so the separation is between corpora and cannot be attributed to framing rather than provenance; this confound is stated before the result wherever the result appears. A second scanning session over all 400 items

gives test-retest $\text{ICC} = 0.8725$ and replicates the separation at $\eta^2 = 0.0888$, while 12.8% of individual items reverse direction between sessions, a rate the measured session noise predicts on its own, so the instrument supports category-level claims and no per-item verdict is made.

The physics is a mean-field treatment in which media enters as a field $h = \alpha X$ on the measured observable. Calibrating α against the corpus does not identify it: the fit performs worse than the sample mean out of sample and the estimate scales with a coupling the study does not measure, so no value is quoted anywhere. What survives is a necessary condition, $\alpha \geq h_c(\beta J)/\Delta X$. The measured spread $\Delta X = 0.1241$ gives $\alpha \geq 4.2938$ at $\beta J = 2$, a constraint on any future proposal of this mechanism that depends on the range of the observable rather than on why that range arises.

Publications arising from this work

Three papers were written from this project, and co-authorship with the supervisor is under discussion. Their status at the time of submission is given as it stands, and no paper is described as published unless it has been accepted.

1. **Bounding the external field in mean-field opinion dynamics: what any content observable must satisfy to drive a transition.** To be submitted to the *International Journal of Modern Physics C*. Corresponds to Chapter 3, and states the bound $\alpha \geq h_c(\beta J) / \Delta X$ together with the screening criterion that follows from it. *Status: complete, awaiting submission.*
2. **A cortical-proxy observable for emotionally manipulative media: an instrument, a 400-item corpus, and the field bound it yields.** To be submitted to the *Journal of Computational Social Science*, as the companion to the first. Corresponds to Chapters 4 to 6, and supplies the measured spread that the first paper's inequality consumes. *Status: complete, awaiting submission.*
3. **What does a released average-subject brain encoder actually predict? A measured noise ceiling and a replication protocol.** To be submitted to *Neurons, Behavior, Data analysis, and Theory*. Reports the inter-subject noise ceiling measured from public data, and the two silent defects behind a value this project withdrew. *Status: drafted; the prediction pass over the first 120 s is reported, and the full-episode pass remains.*

The instrument, the corpus, the analysis scripts and the artifacts every reported number is read from are public at <https://github.com/brn-mwai/monarch>, and the working instrument runs at <https://monarch-4iy.pages.dev>.

Acknowledgements

I thank my supervisor, Dr. Songa Mutambi, for approving a change of instrument mid-project when the released checkpoint turned out not to see the structure the proposal named, and for insisting that the reformulated index be defended on its own terms rather than presented as a minor substitution.

This work runs on public artifacts. TRIBE v2 and its released checkpoint come from Meta AI Research; the corpus is assembled from ISOT, SemEval-2019 Task 4, Webis-Clickbait-17 and PubMed; the fMRI recordings used to measure the noise ceiling come from the Algonauts 2025 release. The GPU time came from Kaggle’s free tier, and the constraints of that tier shaped how the corpus was scanned.

Contents

Declaration	iii
Abstract	v
Publications arising from this work	vii
Acknowledgements	ix
List of Figures	xvi
List of Tables	xviii
Abbreviations and Symbols	xix
1 Introduction	1
1.1 The problem	1
1.2 The amendment, stated at the outset	1
1.3 Research questions	2
1.4 Objectives	2
1.5 Contribution	3
1.6 The longer aim, and what this thesis is	3
1.7 Structure of this document	5
2 Literature Review	7
2.1 Opinion dynamics as statistical physics	7
2.1.1 How the external field is treated	7
2.1.2 The exception, and what it leaves open	8
2.2 Predictive neural encoding	8
2.2.1 What a released checkpoint does and does not guarantee	9

2.3	The gap	9
3	Theoretical Framework	11
3.1	The mean-field model	11
3.2	The Landau expansion, with coefficients derived	11
3.3	Numerical validation of the solver	12
3.4	The spinodal field	13
3.5	The bound	15
3.6	A screening criterion	15
3.7	Assumptions	17
4	Methodology	19
4.1	The apparatus	19
4.1.1	The instrument as published	19
4.1.2	Checkpoint facts, taken from the artifact	21
4.2	The observable	21
4.3	The corpus	21
4.4	Execution	23
4.5	Analysis	23
4.5.1	Power, computed before the analysis	24
4.5.2	Rules fixed before the data existed	24
4.6	Reproducibility	24
5	Results	27
5.1	The corpus as scanned	27
5.2	RQ I: the instrument and its classifier performance	28
5.2.1	The ratio form is undefined on real content	28
5.2.2	Classifier performance	28
5.2.3	What a bag of words does on the same label	29
5.3	RQ II: distribution across categories	30
5.3.1	The two components, taken separately	31
5.3.2	The confound, stated before the interpretation	34

5.4	RQ III: phase structure and the coupling	34
5.4.1	The coupling is unidentified	35
5.4.2	The bound	36
5.5	Statistical power	36
5.6	RQ IV: out of scope	37
5.7	Summary of the answers	37
6	Discussion	39
6.1	What the separation can and cannot support	39
6.1.1	The evidence that argues against a pure source artifact . . .	39
6.1.2	The experiment that would decide it	40
6.2	Why high outrage does not separate	41
6.3	The instrument's standing	42
6.3.1	The checkpoint problem	42
6.3.2	What this means for building on foundation models	43
6.4	The coupling, and what the bound does instead	44
6.5	Session reliability, and what it permits	44
6.6	Limitations	47
6.7	What would be required next	49
6.8	Screening media designed to hold attention	50
7	Conclusion	53
7.1	What was promised	53
7.2	What was delivered	53
7.2.1	The apparatus	53
7.2.2	The framework	54
7.2.3	The measurement	54
7.3	What is claimed, and what is not	55
7.4	Contribution	55
7.5	Future work	56
7.6	Code, data and the working instrument	56
7.7	Closing	56

List of Figures

3.1	Free energy (3.3) at zero field with the derived coefficients (3.5), at $\beta J = 0.8, 1.0$ and 1.4 . The second minimum appears only above the critical coupling.	13
3.2	Order parameter across the transition. The fitted exponent is 0.4872 against an exact $1/2$	14
3.3	Phase diagram in $(\beta J, h)$. The shaded region is where two opinion states coexist; its upper boundary is the spinodal $h_c(\beta J)$	14
3.4	Coupling required before a content observable of spread ΔX can drive the population across the spinodal.	16
4.1	The published corpus page	20
4.2	A per-vertex map on the published page	20
4.3	The two ROI unions on the fsaverage5 surface: affective-salience in orange (1,030 vertices) and deliberative-control in blue (851). This figure is a definition and carries no data.	22
5.1	Distribution of the signed index by category over all 400 items. . .	31
5.2	Every item ranked by the signed index, coloured by category. . . .	32
5.3	A_{aff} , A_{del} and their signed difference by category over all 400 items. High outrage is elevated on both components and flat on the difference.	33
5.4	Predicted response for one fear-activating item at every one of the 20,484 cortical vertices, on the pial surface. Vertices below the map's 70th percentile are left unpainted so the anatomy remains readable; the threshold is stated in the figure. This is the encoder's own output and not a smoothed or interpolated field.	35

5.5	Free energy at each category's measured mean, swept over α . The coupling is swept and never fitted, and no value of $\hat{\alpha}$ is quoted. . .	36
-----	---	----

List of Tables

4.1	Corpus composition. Each category is drawn from exactly one source dataset, which is the confound discussed throughout: category and provenance cannot be separated.	22
5.1	Signed index by category over all 400 scanned items. Every category mean is negative, so deliberative activation leads throughout. . . .	27
5.2	Classifier performance against the pre-scan manipulative label. AUC is the headline figure; the threshold behind F1 was fitted in sample and F1 is labelled accordingly.	28
5.3	Baselines against the same manipulative label. The lower block restricts the corpus to the 150 items drawn from the ISOT collection, the only origin corpus contributing both classes.	29
5.4	Welch comparisons of each category against neutral informational. The separation is carried by fear-activating content; high outrage does not separate, and it is the category the proposal expected to separate most.	31
5.5	Category means of each component and of their difference, over all 400 items.	32
5.6	Welch comparisons of each component against neutral informational. High outrage moves both components by nearly the same amount and the difference between them by nothing; fear-activating content moves one component only.	33
5.7	The field bound at three values of the reduced coupling. The measured spread $\Delta X = 0.1241$ converts each critical field into the coupling any proposal of this mechanism would require.	36

5.8	Smallest effects this design could detect at $\alpha = 0.05$ and power 0.80, fixed before the corpus was analysed.	37
6.1	Agreement between the two scanning sessions over all 400 items. ICC is reported beside r because a session that shifted every score by a constant would still correlate perfectly.	45

Abbreviations and Symbols

NAA	Network Activation Asymmetry, the observable this work measures
A_{aff}	Mean predicted activation across affective-salience regions
A_{del}	Mean predicted activation across deliberative-control regions
$\text{NAA}_{\text{signed}}$	$A_{\text{aff}} - A_{\text{del}}$, the signed index used throughout
ΔX	Measured spread of the observable across the corpus
α	Coupling between the observable and the model's external field, $h = \alpha X$
βJ	Reduced interaction coupling in the mean-field model
h_c	Critical field at which the mean-field free energy loses bistability
η^2	Proportion of variance explained, the effect size for category separation
AUC	Area under the receiver operating characteristic curve
ICC	Intraclass correlation coefficient, absolute-agreement form ICC(2,1)
fMRI	Functional magnetic resonance imaging
ROI	Region of interest
TR	Repetition time, one fMRI sampling interval

Introduction

1.1 THE PROBLEM

Two articles can report the same event and leave a reader in entirely different states. One informs; the other provokes. Readers notice the difference and cannot quantify it, and the platforms that distribute both have no objective measure of which they are amplifying. The difference is real, consequential and unmeasured.

Existing approaches measure the text. Sentiment analysis counts affective vocabulary; readability scores count syllables and clause length; stance detection classifies a position. None of them measures what the content is predicted to *do* to the person reading it.

This project takes a different route. Rather than characterising the words, it estimates the response: a released neural encoder predicts cortical activity from the content, and the measurement is taken on that prediction. The instrument is a thermometer for one specific property of media, and the thesis is an account of building it, running it, and stating carefully what its readings support.

1.2 THE AMENDMENT, STATED AT THE OUTSET

Three commitments in the original proposal could not be met as written. They are set out here rather than disclosed in the results chapter, because a reader who meets them for the first time beside a result will reasonably suspect they were discovered to explain it. Each was identified before the measurement was taken, and each is documented in the amendment approved by the supervisor.

First, the observable is cortical. The proposal framed the index around sub-cortical affective structures. The released checkpoint predicts cortical surface

activity only, over 20,484 vertices, and contains no subcortical output. The index is therefore a *cortical proxy*: a contrast between two cortical ROI unions, one associated with affective salience and one with deliberative control. No claim about subcortical structures appears anywhere in this document.

Second, the index is a difference, not a ratio. Equation (4) of the proposal defines the index as $A_{\text{aff}} / (A_{\text{del}} + \delta)$. The encoder emits standardised activation, so an ROI mean is negative about as often as positive and the ratio sign-flips or diverges. On the completed corpus it is undefined for 17.25% of items. The measurement uses the signed difference $A_{\text{aff}} - A_{\text{del}}$, which is defined everywhere and is the form the physics requires.

Third, the corpus is text-only. The proposal specified 1,500 multimodal items. The delivered corpus is 400 text items in four balanced categories, and Research Question IV, which required audio and video, is out of scope.

1.3 RESEARCH QUESTIONS

RQ I Can predicted cortical activation serve as a quantitative signature of emotionally manipulative media content, and how does the resulting index perform as a classifier against pre-scan dataset-derived labels?

RQ II How is the index distributed across content categories, and do the categories separate?

RQ III What phase structure does a mean-field opinion model exhibit under a field derived from this observable, and what does the measured spread imply for the coupling?

RQ IV How do text, audio and video contribute to the index? *Out of scope under the amendment.*

1.4 OBJECTIVES

- (i) Build a labelled corpus of media items in balanced categories.
- (ii) Run the encoder over the corpus and compute the index for every item.

- (iii) Characterise the distribution of the index.
- (iv) Calibrate the coupling between the index and the model's field.
- (v) Derive the phase structure of the opinion model.
- (vi) Validate the index as a classifier against pre-scan dataset-derived labels.
- (vii) Deliver the measurement apparatus as working open-source code.

Objective (vii) is the primary deliverable. The proposal states this explicitly and it governs how the work should be assessed: the apparatus is the promise, and the empirical result is what the apparatus returned when it was run.

1.5 CONTRIBUTION

Three, developed across the chapters that follow.

A **bound** on the coupling between any content observable and an opinion-dynamics field, requiring only the observable's measured spread and the model's own spinodal, and evaluable before data are collected (Chapter 3).

An **instrument**, delivered as open-source code, together with the documented finding that the released checkpoint is cortical-only and that its agreement with real cortex is contested (Chapter 4).

A **measurement** over 400 items, reported with the design's sensitivity attached and its central confound stated before its interpretation (Chapters 5 and 6).

1.6 THE LONGER AIM, AND WHAT THIS THESIS IS

The proposal defence framed a use for an instrument of this kind that is larger than any single study: a tool that ordinary people, rather than researchers, could use to ask whether a piece of media is built to manipulate, whether a programme is built to keep a child watching, or how a person is likely to react to something before they see it. That aim is worth stating because it explains why this work was undertaken, and it is worth separating cleanly from what has been delivered, because the distance between the two is the honest part.

This thesis is the first phase of that programme, and only the first. What it delivers is an instrument that computes one number per item from content, a corpus on which that number's behaviour is characterised, an account of where the measurement is reliable and where it is not, and a bound that constrains what any successor may assume. What it does not deliver is a tool anyone should use to judge a particular piece of media.

The gap is not a matter of polish. It is four specific things, each one a limitation measured in this document rather than an estimate:

1. Whether the instrument measures anything about brains is unresolved.

The encoder's agreement with real cortex is contested and the replication is unfinished (§6.3). Until that is settled every reading is a statement about a model's output.

2. The measurement cannot yet be attributed to content. Category and source dataset coincide in this corpus, so what separates could be provenance rather than framing (§5.3.2). A crossed corpus fixes this and needs no new instrumentation.

3. Individual items cannot be judged. 12.8% of items reverse direction between scanning sessions (§6.5), so group comparisons are supportable and a verdict on one article, or one programme, is not.

4. Nothing here connects the number to behaviour. Predicting how a person will react, or whether they will keep watching, or what they will buy, is a claim about behaviour. This instrument predicts cortical response, and no behavioural data was collected against which that prediction could be checked.

Stated as phases, the first is done and the rest are not:

Phase	What it establishes	Status
1	An instrument, a measured observable, its reliability, its confound, and the bound	This thesis
2	Whether the encoder tracks real cortex; a corpus that separates framing from source	Specified, partly run
3	Video rather than text, the observable over time rather than pooled, validated against viewing behaviour	Designed, not run
4	Response for a particular person rather than an average	Out of reach here

Phase 4 deserves its own sentence, because it is the one most easily promised. The released checkpoint is loaded in an average-subject configuration that its own loader forces (§6.3), so it has no representation of an individual to predict for. Reaching individual response would require per-subject models, individual imaging data and the consent framework that comes with it. It is not a refinement of what is built here; it is a different study.

None of this reduces what the first phase settles. A bound that any future proposal must satisfy does not expire when the instrument improves, and a corpus with its reliability characterised is more useful to a successor than a larger one without. The claim made in this document is exactly that: a first phase, reported with its limits measured.

1.7 STRUCTURE OF THIS DOCUMENT

Chapter 2 reviews the literature on opinion dynamics as statistical physics and on predictive neural encoding, and identifies the gap this work addresses. Chapter 3 derives the theoretical framework and the bound. Chapter 4 describes the instrument, the corpus and the analysis. Chapter 5 reports the results. Chapter 6 discusses what they support. Chapter 7 concludes.

Literature Review

This work sits between two literatures that rarely meet: the statistical physics of opinion dynamics, which supplies the model, and predictive neural encoding, which supplies the observable. The gap it addresses lies precisely at their junction.

2.1 OPINION DYNAMICS AS STATISTICAL PHYSICS

The programme of treating collective opinion with the tools of statistical mechanics is established and productive. Castellano, Fortunato and Loreto’s review [1] remains the standard survey: binary attitudes map to Ising spins, social conformity to a ferromagnetic coupling J , and external influence to a field h . Galam’s review of his own models [2] traces the same lineage through majority-rule dynamics, and Sznajd-Weron and Sznajd [3] contribute the outward-influence variant that carries their name. Stauffer [4] gives an accessible account of why two-dimensional Ising models apply to social questions at all.

The recent review by Mullick and Sen [5] surveys a century of Ising-inspired sociophysics across opinion dynamics, markets, segregation and language. Notably for this work, it lists data-driven calibration against large-scale social datasets among the field’s *open directions* rather than among its accomplishments.

2.1.1 How the external field is treated

Where these models represent media explicitly, the field is assigned rather than measured. Crokidakis [6] introduces mass media into the two-dimensional Sznajd model as a probability p of following the media opinion, sweeps it, and finds the phase transition absent above $p \approx 0.18$. Azhari and Muslim [7] apply an external field to majority-rule and q -voter dynamics on the complete graph and sweep its

magnitude. Freitas, Vieira and Anteneodo [8] study imperfect bifurcations and cusp catastrophes under external bias, parameterised by sensitivity and direction rather than by any measured quantity. The review of models with external effects by Sîrbu and colleagues [9] surveys this class without any of them measuring the external term.

In each case the swept parameter is the quantity whose real-world magnitude is at issue. The literature can therefore demonstrate that a population tips at some h^* without establishing that any real influence attains it.

2.1.2 The exception, and what it leaves open

One recent study breaks the pattern. Korbel, Dahdoul and Thurner [10] calibrate a double random-field Ising model of elections directly to campaign expenditure and extract a critical spending threshold from historical results, reporting it as the first time thermodynamic parameters and a critical threshold have been taken from election data in this way.

Two things remain open after it, and they define the gap this work addresses. Their calibration is fitted from outcome data after the fact, so it cannot screen a proposed mechanism before data exist. And it is specific to expenditure, whereas observables proposed for media *content* remain uncalibrated. What is missing in both cases is a necessary condition stated on the coupling itself, which is what Chapter 3 supplies.

2.2 PREDICTIVE NEURAL ENCODING

The second literature concerns models that predict brain responses to naturalistic stimuli. The relevant instrument here is TRIBE [11], a trimodal encoder that combines pretrained text, audio and video representations to predict fMRI responses to film, trained on the Courtois NeuroMod recordings and evaluated over 1,000 cortical parcels.

Such models make a measurement possible that was not before: rather than characterising a stimulus by its surface features, one can estimate the response it is predicted to elicit. That is the move this project makes.

2.2.1 What a released checkpoint does and does not guarantee

A technical report describes an architecture and a training procedure. The weights released under its name are a separate artifact, and their properties must be established from the artifact itself.

Three such properties matter here and none is visible from the publication. The checkpoint used in this work predicts cortical surface activity only. Its loader forces an average-subject configuration at load time, so per-subject variation is unavailable. And its agreement with real cortex in that configuration is contested: a commercially published audit reports vertex-level anti-correlation against a measured inter-subject ceiling, across four subjects and all seven Yeo networks. That audit is self-published by a company selling the per-subject alternative, uses one stimulus, and reports small magnitudes in either direction; it has not been independently replicated. It should be neither accepted uncritically nor dismissed.

2.3 THE GAP

Placing the two literatures beside one another gives the gap directly.

The sociophysics literature has a field it cannot measure. The encoding literature has a predicted response that has not been used as a field. Neither has stated what a candidate observable would have to satisfy for the mechanism to operate at all.

This work supplies that condition, builds an instrument that produces a candidate observable, measures its spread on a controlled corpus, and reports what follows — including the respects in which the design cannot support the interpretation one would prefer.

Theoretical Framework

3.1 THE MEAN-FIELD MODEL

Take N agents holding binary opinions $s_i = \pm 1$, coupled all-to-all with strength J/N , in an external field h :

$$\mathcal{H} = -\frac{J}{N} \sum_{i<j} s_i s_j - h \sum_i s_i. \quad (3.1)$$

The collective polarisation is $m = N^{-1} \sum_i s_i$, and in the thermodynamic limit it satisfies the self-consistency condition

$$m = \tanh(\beta J m + h), \quad (3.2)$$

with $\beta = 1/k_B T$ and $k_B = 1$ throughout. The field h absorbs the factor β , so it is dimensionless and directly comparable with βJ .

The social reading of β is an inverse noise level. Large β describes a population that copies its neighbours reliably; small β one that decides close to at random. Equation (3.2) has a single root for $\beta J < 1$ at zero field and three for $\beta J > 1$, of which the outer two are stable.

3.2 THE LANDAU EXPANSION, WITH COEFFICIENTS DERIVED

Write the free energy as a quartic in the order parameter,

$$F(m) = a m^2 + b m^4 - h m. \quad (3.3)$$

The coefficients are not free. Requiring that $\partial F/\partial m = 0$ reproduce (3.2) to cubic order fixes them. Inverting (3.2) gives $\operatorname{artanh}(m) = \beta J m + h$, and expanding $\operatorname{artanh}(m) = m + m^3/3 + O(m^5)$ yields

$$(1 - \beta J) m + \frac{1}{3} m^3 - h = 0. \quad (3.4)$$

Matching against $\partial F/\partial m = 2am + 4bm^3 - h$ gives

$$a = \frac{1 - \beta J}{2}, \quad b = \frac{1}{12}. \quad (3.5)$$

This derivation is stated in full because coefficient forms that do not follow from the accompanying self-consistency condition appear in the applied literature, and the symptom is subtle: the free-energy curve retains the correct qualitative double well, so the figure looks right, while its minimum sits away from the m^* the same work's root-finder reports. The double well opens when $a < 0$, that is when $\beta J > 1$, recovering the transition already implied by (3.2).

The susceptibility follows as

$$\chi = \frac{\beta \operatorname{sech}^2(\beta J m^* + h)}{1 - \beta J \operatorname{sech}^2(\beta J m^* + h)}, \quad (3.6)$$

which diverges where the denominator vanishes and is undefined beyond it. It is reported only where the denominator is positive.

3.3 NUMERICAL VALIDATION OF THE SOLVER

Every figure in this framework is drawn with one solver, so the solver is checked first against results known exactly. Mean-field theory fixes three critical exponents: $m^* \sim (\beta J - 1)^\beta$ with $\beta = 1/2$, $\chi \sim |1 - \beta J|^{-\gamma}$ with $\gamma = 1$, and $m \sim h^{1/\delta}$ on the critical isotherm with $\delta = 3$ [13].

One methodological point matters. Solving (3.2) by direct iteration of $m_{k+1} = \tanh(\beta J m_k + h)$ is adequate away from criticality but degrades exactly at $\beta J = 1$, where the map's derivative reaches unity. The exponents are defined in that limit,

so fitting them to iterated values measures truncation error rather than physics: doing so returns $\delta \approx 1.35$ against an exact 3. The exponents are therefore fitted to roots obtained by bracketing, and the iterative solver is checked against the bracketed one away from criticality, where the largest absolute discrepancy over the sampled grid is 1.15×10^{-9} .

Fitting within a window of 0.05 in $|\beta J - 1|$ at 60 points recovers

$$\beta = 0.4872, \quad \gamma = 1.0000, \quad \delta = 3.0042, \quad (3.7)$$

with absolute errors 0.0128, 0.0000 and 0.0042. The residual in β is a finite-window effect: the power law is asymptotic and the quartic term contaminates the slope as the window widens. The run fails if any exponent departs from its exact value by more than a stated tolerance, so a solver regression stops the analysis rather than producing plausible curves.




Figure 3.1: Free energy (3.3) at zero field with the derived coefficients (3.5), at $\beta J = 0.8, 1.0$ and 1.4 . The second minimum appears only above the critical coupling.

3.4 THE SPINODAL FIELD

Above $\beta J = 1$ the free energy has two minima at zero field. Raising the field deepens one and shallows the other; at a finite field the shallower minimum




Figure 3.2: Order parameter across the transition. The fitted exponent is 0.4872 against an exact $1/2$.

disappears. That value is the spinodal field $h_c(\beta J)$, and it is the field a population in the metastable state must experience before it is obliged to move.

It is located by bisection on the presence of a second stationary point, which keeps it consistent with the solver used elsewhere rather than importing a closed form derived under different approximations. It vanishes as $\beta J \rightarrow 1^+$ and grows monotonically, reaching 0.53284 at $\beta J = 2$.




Figure 3.3: Phase diagram in $(\beta J, h)$. The shaded region is where two opinion states coexist; its upper boundary is the spinodal $h_c(\beta J)$.

3.5 THE BOUND

Let an observable X be proposed as the field through a linear map

$$h = \alpha X, \quad (3.8)$$

with α an unknown coupling carrying the units. Any real observable is bounded: across the content a population plausibly encounters, X spans a finite range ΔX , which is an empirical property of the observable and the corpus rather than of the model. The largest field the mechanism can deliver is therefore $\alpha \Delta X$, and requiring it to reach the spinodal gives

$$\alpha \Delta X \geq h_c(\beta J) \quad \Longleftrightarrow \quad \alpha \geq \frac{h_c(\beta J)}{\Delta X}. \quad (3.9)$$

Three properties of (3.9) carry the argument of this thesis.

It is **necessary and not sufficient**. Clearing it does not establish that content tips populations; falling below it excludes the mechanism within the model's assumptions.

It **requires no fit**. ΔX is obtained by measuring the observable across a corpus, which is far cheaper than estimating a coupling to collective behaviour, and h_c comes from the model.

It **converts a null into a constraint**. A study that estimates α and obtains an interval containing zero has, alone, only the statement that no coupling was detected. Combined with (3.9), the same interval says the coupling would have to exceed $h_c(\beta J)/\Delta X$ for the mechanism to operate, and that the measurement excludes that range.

3.6 A SCREENING CRITERION

The bound supports a procedure that runs before data collection.

1. State the observable and the map $h = \alpha X$ explicitly, with units.
2. Measure or bound ΔX on a modest sample. This needs no outcome variable.




Figure 3.4: Coupling required before a content observable of spread ΔX can drive the population across the spinodal.

3. Choose the range of βJ the population plausibly occupies and read $h_c(\beta J)$.
4. Evaluate $\alpha_{\min} = h_c(\beta J)/\Delta X$ and ask whether a coupling of that size is defensible on independent grounds.
5. If it is not, the mechanism is excluded within the model and the expensive measurement is unnecessary. If it is, proceed, and report ΔX alongside the estimate of α so the bound can be applied to the result.

Step 4 is the weakest link, since α has no independent scale in most applications. The criterion remains informative in relative terms: two candidate observables can be ranked by $\alpha_{\min}$, and an observable whose spread demands an implausibly large coupling is a poor choice of field regardless of how well it correlates with anything.

That ranking requires one normalisation to be meaningful. Since $h_c(\beta J)$ is fixed by the population and does not depend on the observable, $\alpha_{\min} = h_c/\Delta X$ orders candidates purely by $1/\Delta X$, and ΔX carries the units of X . Two observables measured on different scales therefore cannot be compared as they stand: rescaling an observable by a constant rescales $\alpha_{\min}$ by its inverse and can reorder the ranking arbitrarily. The comparison is defined only once every candidate is expressed on a common scale, and the natural choice is to standardise each observable by

its own population standard deviation before the spread is taken, so that ΔX is a spread in standard deviations and $\alpha_{\min}$ is a required coupling per standard deviation of the observable. This thesis evaluates one observable and so performs no ranking; the $\Delta X = 0.1241$ carried through Chapter 5 is a raw spread in the encoder's standardised-activation units, which is 5.83 standard deviations of the index across the corpus. Either scale is a legitimate report, and the bound is unaffected as long as α is read in the matching units. What is not legitimate is placing a second observable beside it without first putting both on the same scale: that comparison is not a comparison but a statement about units.

3.7 ASSUMPTIONS

The framework assumes all-to-all coupling, a single homogeneous population, a static field and binary opinions. Network structure and heterogeneous susceptibility generally lower h_c , so (3.9) is a bound within this model rather than a universal one. The linear map $h = \alpha X$ is the simplest possible coupling and is itself an assumption; a saturating map would change the relation between ΔX and the maximum attainable field. Finally, the spinodal is the threshold relevant to a population trapped in a metastable state; a population at coexistence can be moved by a smaller field given enough time, so the bound demands more of α than a nucleation argument would.

Methodology

4.1 THE APPARATUS

The instrument takes a passage of text and returns a predicted cortical response. Five stages run in sequence.

1. **Speech synthesis.** The passage is rendered to audio, since the encoder expects naturalistic audiovisual input rather than written text.
2. **Word timings.** A forced-alignment transcriber recovers per-word timings, which the encoder requires to align the text stream to the audio.
3. **Embedding.** Text and audio are embedded by the pretrained models the encoder was built on.
4. **Prediction.** The encoder, TRIBE v2 [12], returns activation over 20,484 fsaverage5 cortical surface vertices per time point.
5. **Reduction.** The prediction is mean-pooled over time and reduced to two ROI means, from which the index is computed.

4.1.1 The instrument as published

The apparatus is not described only in this chapter. It runs publicly at <https://monarch-4iy.pages.dev>, and the corpus page there reads the same measurement file this dissertation reports from, so the figures a reader sees on the page are the figures in Chapter 5 rather than a summary written about them.




Figure 4.1: The corpus page of the instrument at <https://monarch-4iy.pages.dev/corpus/>, on 2 September 2026. The four values it reports are the ones this dissertation reports: 400 items scanned, 331 of 400 with the ratio form defined, a spread of $\Delta X = 0.1241$, and a range of -0.070 to $+0.054$. The page reads `data/corpus.json`, exported from the same `corpus_naa.csv` the analysis scripts read, so the two cannot drift apart.




Figure 4.2: One item's predicted response on the same page, painted at every one of the 20,484 vertices and thresholded at $+0.1043$, with the two region means and their difference printed beneath it. The map is the encoder's stored output for that item; the page renders no activation it did not scan, and shows an explicit failure notice rather than a substitute when a prediction is unavailable.

4.1.2 Checkpoint facts, taken from the artifact

Three properties of the released checkpoint were established by inspecting the artifact rather than by reading the accompanying publication, using `scripts/verify_tribe_checkpoint.py`. Its configuration declares `mesh:fsaverage5`, and the smoke test returns arrays of shape $(T, 20484)$, which is 2×10242 vertices: the output is cortical-only. Its loader sets `average_subjects = True` at load time, so the subject term is unavailable. Both facts are recorded in the run log with the file and line that establish them.

4.2 THE OBSERVABLE

Two ROI unions are defined over the cortical surface: an affective-salience set of 1,030 vertices and a deliberative-control set of 851, with no overlap. For an item’s mean-pooled prediction, A_{aff} and A_{del} are the mean activation over each set.

The index is the signed difference

$$\text{NAA}_{\text{signed}} = A_{\text{aff}} - A_{\text{del}}, \quad (4.1)$$

positive when the affective-salience set is predicted to lead. The ratio form is computed and stored alongside it, and the rate at which it is undefined is reported as a result in its own right (§5.2).

The observable is a **cortical proxy**. It is not a measurement of any person, and the regions it contrasts are cortical unions rather than the subcortical structures the original proposal named.

4.3 THE CORPUS

400 text items in four categories of 100: fear-activating, high outrage, reward-hook, and neutral informational.

Items are length-matched: category word counts are 167.4, 163.5, 163.7 and 162.2 with standard deviations near 11. Each item carries the labels it held before scanning — `manipulative`, `credibility`, `partisan_intensity` and its source —




Figure 4.3: The two ROI unions on the fsaverage5 surface: affective-salience in orange (1,030 vertices) and deliberative-control in blue (851). This figure is a definition and carries no data.

Category	Source	n
Fear activating	ISOT-fake	100
High outrage	SemEval-2019 Task 4	100
Reward hook	Webis-Clickbait-17	100
Neutral informational	PubMed, ISOT-true	50 + 50

Table 4.1: Corpus composition. Each category is drawn from exactly one source dataset, which is the confound discussed throughout: category and provenance cannot be separated.

and these are the only ground truth in the dataset, since the predicted activations have none. The sources are the ISOT fake and true news collections [14], SemEval-2019 Task 4 [15], Webis-Clickbait-17 [16], and PubMed abstracts retrieved through the NCBI utilities [17]. The manipulative label is derived from the source dataset rather than annotated per item.

NELA-GT-2021, named in the proposal, was withdrawn by its authors in January 2024. SemEval-2019 Task 4 is substituted; it is openly licensed and labels the article rather than the publisher.

Each category is drawn from a single source. Category and provenance are therefore confounded by construction. This is a property of the design, is stated here rather than discovered later, and bounds the interpretation throughout Chapters 5 and 6.

4.4 EXECUTION

The scan is the only step requiring a GPU and the only one that cannot be repeated cheaply. It was executed in four sessions on a Tesla P100 at a measured 64 to 78 seconds per item. Every item is appended and flushed as it completes, so an interrupted session retains everything computed, and a resumed run skips items already present.

Output was verified before any analysis. The verification script, `scripts/verify_scan_output`, checks every scanned row against the source corpus on text, category and identifier, requires distinct values, and confirms that the stored difference agrees with the two stored means to the precision at which they are written. All 400 rows passed.

4.5 ANALYSIS

RQ I scores the index against the manipulative label by AUC, with F1, precision and recall reported alongside and labelled as fitted in sample.

RQ II tests separation across the four categories by one-way ANOVA, reporting F , p and η^2 , with per-category descriptives and Welch contrasts against the neutral baseline.

RQ III evaluates the free energy and susceptibility across a range of α rather than at a fitted value, and applies the bound of §3.5 to the measured spread.

4.5.1 Power, computed before the analysis

The design's sensitivity was computed from the design alone, before the corpus was analysed, using noncentral F for the ANOVA and a seeded Mann-Whitney simulation for the AUC. At $\alpha = 0.05$ and 80% power, the smallest detectable effects are $\eta^2 = 0.0268$ and $\text{AUC} = 0.5916$.

Computing this in advance is deliberate. A power statement produced after a result is a rationalisation; produced beforehand it states what the design could and could not have found, and it is reported alongside every result whichever way they came out.

4.5.2 Rules fixed before the data existed

Five, recorded in the project's post-scan runbook before the corpus was complete.

AUC is the headline and F1 is not, since F1's threshold is fitted in sample. An AUC below 0.5 would be reported as measured rather than sign-flipped. An AUC at or above 0.95 would be treated as a leak until its source was identified. A non-significant result would be reported as a null with the excluded effect size attached, not as an absence of effect. And no value of $\hat{\alpha}$ would be quoted anywhere, whatever the calibration returned.

4.6 REPRODUCIBILITY

Every number in this thesis is produced by a script in the repository and regenerates from the corpus file. The instrument, the analysis, the power calculation and the physics are open source. The run log records each execution with its command,

its output and its failures, including the runs that failed and why, so that the path to the reported numbers is auditable rather than asserted.

Results

This chapter answers the research questions in order. Every number in it is produced by a script in the repository and is reproducible from the corpus file `data/final/corpus_naa.csv`; none is transcribed by hand. Where the design cannot support a claim, the chapter says so at the point the number appears rather than deferring it to the limitations section.

One convention is fixed before any result: no value of $\hat{\alpha}$ appears anywhere in this document. Section 5.4 explains why the coupling is unidentified rather than merely uncertain.

5.1 THE CORPUS AS SCANNED

The corpus is 400 text items in four categories of 100. All 400 were scanned; the output was verified row for row against the source corpus on text, category and identifier before any analysis was run, using `scripts/verify_scan_output.py`.

Category	n	Mean	SD
Fear activating	100	−0.0018	0.0204
High outrage	100	−0.0185	0.0208
Neutral informational	100	−0.0191	0.0162
Reward hook	100	−0.0128	0.0228

Table 5.1: Signed index by category over all 400 scanned items. Every category mean is negative, so deliberative activation leads throughout.

Two features of this table matter before any test is applied. Every category mean is negative, so the deliberative-control network is predicted to respond more strongly than the affective-salience network for all four categories; the categories differ in how far below zero they sit, not in which network leads. And the standard deviations are of the same order as the differences between the means, which is

the reason the separation test in Section 5.3 is reported with an effect size rather than a p -value alone.

5.2 RQ I: THE INSTRUMENT AND ITS CLASSIFIER PERFORMANCE

5.2.1 The ratio form is undefined on real content

Equation (4) of the proposal defines the index as a ratio, $A_{\text{aff}}/(A_{\text{del}} + \delta)$. The encoder emits standardised activation, so an ROI mean sits near zero and is negative about as often as positive, and the ratio then sign-flips or diverges. On the completed corpus the ratio is undefined for **69 of 400 items**, a rate of 17.25%. Those items are counted and reported, never dropped or filled.

The measurement therefore uses the signed difference,

$$\text{NAA}_{\text{signed}} = A_{\text{aff}} - A_{\text{del}}, \quad (5.1)$$

which is defined everywhere and is what the field mapping in Chapter 3 requires: the Landau field needs a signed scalar, not a dimensionless ratio. The classification banding follows the same constraint. The signed form supports a two-way split on sign, and no magnitude threshold has been established from any corpus, so the three-band scheme of the proposal becomes two bands.

5.2.2 Classifier performance

Scored against the manipulative label the items carried before they were ever scanned, over 300 labelled manipulative and 100 labelled neutral:

Metric	Value
AUC (headline)	0.6274
F1 (threshold fitted in sample)	0.6680
Precision at that threshold	0.8426
Recall at that threshold	0.5533

Table 5.2: Classifier performance against the pre-scan manipulative label. AUC is the headline figure; the threshold behind F1 was fitted in sample and F1 is labelled accordingly.

AUC is the headline and F1 is not. The threshold that produces F1 is chosen on the same data it is scored on, which inflates it; AUC is threshold-free and does not reward that choice. A null corpus can return a respectable F1 at an AUC below chance, so quoting F1 alone would present a failure as a success.

The measured AUC of 0.6274 clears 0.5916, the smallest AUC this design could detect at 80% power (Section 5.5), and it lies in the direction the proposal predicted. It is nonetheless a modest separation, and no per-item verdict is derived from it in this work: at this AUC, a correct-or-incorrect label on individual items would overstate what the measurement supports.

5.2.3 What a bag of words does on the same label

An AUC has no meaning in isolation. The first question a reader should ask is whether a lexical model, costing no GPU time at all, does better on the same rows and the same label. It does, and by a wide margin.

Two baselines were run with `scripts/baselines.py`: TF-IDF over unigrams and bigrams with logistic regression, scored on out-of-fold predictions from five-fold stratified cross-validation, and a sentiment score, the VADER compound valence [18], used directly as a score with nothing fitted. Confidence intervals are 2,000-draw bootstrap percentile intervals on the same rows.

Model	Split	n	AUC	95% CI
NAA signed index	none, nothing fitted	400	0.6274	[0.5702, 0.6867]
VADER compound	none, nothing fitted	400	0.5392	[0.4757, 0.6035]
VADER compound	none, nothing fitted	400	0.5300	[0.4693, 0.5923]
TF-IDF + logistic	5-fold stratified	400	0.9758	[0.9601, 0.9879]
NAA signed index	ISOT only	150	0.7290	[0.6431, 0.8132]
VADER compound	ISOT only	150	0.4019	[0.3098, 0.4943]
TF-IDF + logistic	ISOT only, 5-fold	150	0.9724	[0.9483, 0.9901]

Table 5.3: Baselines against the same manipulative label. The lower block restricts the corpus to the 150 items drawn from the ISOT collection, the only origin corpus contributing both classes.

The lexical baseline wins, and the sentiment baseline does not. TF-IDF reaches 0.9758 against the index’s 0.6274, and the two intervals are nowhere near

overlapping. The sentiment score does not clear chance: its interval spans 0.5 on the full corpus, and on the restricted set it lies wholly below 0.5, so there it ranks the label the wrong way round. No claim that the index is a better detector of manipulative content than counting words is made anywhere in this thesis, and this table is here so that the claim cannot be read in by implication.

Two qualifications belong with the number rather than after it. The first is that leave-one-source-out, the obvious correction for the confound of §5.3.2, is not available: every source dataset in the corpus contributes exactly one class, so a held-out source leaves a single-class test fold and no per-fold AUC exists. The restricted block is the available substitute. ISOT-true and ISOT-fake come from one collection, so a model scored on those 150 items cannot win by recognising which corpus a row came from, and TF-IDF stays at 0.9724 there. Its advantage over the index is therefore not only the cross-corpus confound.

The second is that both TF-IDF figures sit above 0.95, the level at which this project’s own validation module flags a result as more likely a leak than a discovery (`app/services/validation.py`). The ISOT true/fake split is known to carry a wire-service dateline artifact, which the corpus builder strips and verifies, and a near-ceiling lexical result on a 150-item split is the signature that some residual stylistic marker survives it. A probe trained on the same features to recover `source_dataset` itself, rather than the label, reaches 66.25% accuracy against a 25% chance rate, which confirms the vocabulary is substantially source-diagnostic.

What the comparison establishes is therefore narrow, and it is the honest reading: on this corpus a lexical model separates the label far better than the index does, and a good part of what it separates is provenance. The index is carried forward here as the field observable Chapter 3 requires, not as a competitive classifier.

5.3 RQ II: DISTRIBUTION ACROSS CATEGORIES

A one-way analysis of variance across the four categories gives

$$F = 15.779, \quad p = 1.035 \times 10^{-9}, \quad \eta^2 = 0.1068. \quad (5.2)$$

The effect is roughly four times the $\eta^2 = 0.0268$ the design was powered to detect, and power at the observed value is 1.0. The categories separate, and this is not the null that both partial runs pointed toward.

The separation is not evenly distributed. Welch tests of each category against the neutral baseline give:

Category vs neutral	t	p	Cohen's d
Fear activating	+6.639	3.283×10^{-10}	+0.939
Reward hook	+2.254	2.539×10^{-2}	+0.319
High outrage	+0.212	8.320×10^{-1}	+0.030

Table 5.4: Welch comparisons of each category against neutral informational. The separation is carried by fear-activating content; high outrage does not separate, and it is the category the proposal expected to separate most.

High outrage does not separate from neutral at all. That is the category the proposal expected to separate most strongly. Fear-activating content carries the result almost entirely, with reward-hook content contributing a small effect.




Figure 5.1: Distribution of the signed index by category over all 400 items.

5.3.1 The two components, taken separately

The signed index is a difference, and a difference discards information the corpus already holds. A_{aff} and A_{del} are recorded per item, so each can be compared




Figure 5.2: Every item ranked by the signed index, coloured by category.

against the neutral baseline on its own. This matters for high outrage in particular, where a large symmetric response and no response are indistinguishable in a difference.

Category	A_{aff}	A_{del}	signed
Fear activating	0.03206	0.03385	-0.00179
High outrage	0.02869	0.04719	-0.01850
Neutral informational	0.01699	0.03605	-0.01906
Reward hook	0.02310	0.03585	-0.01275

Table 5.5: Category means of each component and of their difference, over all 400 items.

High outrage is not a null in the components. It raises A_{aff} by $d = +0.507$ and A_{del} by $d = +0.491$, both at $p < 10^{-3}$, and it moves the difference between them by $d = +0.030$ at $p = 0.832$. The two effects sit within four hundredths of a standard deviation of each other. The category that fails to separate on the index is the category that moves both networks together.

Fear-activating content behaves differently in exactly the way that reading predicts. It raises A_{aff} by $d = +0.613$ and leaves A_{del} where it was ($d = -0.098$, $p = 0.488$), so the whole of its $d = +0.939$ on the index is a one-sided movement.

Category vs neutral	Component	t	p	Cohen's d
Fear activating	A_{aff}	+4.333	2.365×10^{-5}	+0.613
Fear activating	A_{del}	-0.696	4.875×10^{-1}	-0.098
Fear activating	signed	+6.639	3.283×10^{-10}	+0.939
High outrage	A_{aff}	+3.584	4.263×10^{-4}	+0.507
High outrage	A_{del}	+3.468	6.521×10^{-4}	+0.491
High outrage	signed	+0.212	8.320×10^{-1}	+0.030
Reward hook	A_{aff}	+1.842	6.691×10^{-2}	+0.261
Reward hook	A_{del}	-0.063	9.501×10^{-1}	-0.009
Reward hook	signed	+2.254	2.539×10^{-2}	+0.319

Table 5.6: Welch comparisons of each component against neutral informational. High outrage moves both components by nearly the same amount and the difference between them by nothing; fear-activating content moves one component only.

Reward-hook content sits between the two, with a small one-sided lift that does not reach significance on its component alone.

The pattern holds at the level of the whole design. A one-way analysis of variance on each component separately gives $F = 7.451$, $p = 7.276 \times 10^{-5}$ for A_{aff} and $F = 6.399$, $p = 3.052 \times 10^{-4}$ for A_{del} , against $F = 15.779$, $p = 1.035 \times 10^{-9}$ for the difference. Both components vary by category; the difference varies more, because for two of the four categories one component moves while the other holds still.




Figure 5.3: A_{aff} , A_{del} and their signed difference by category over all 400 items. High outrage is elevated on both components and flat on the difference.

This does not resolve the confound below, and it says nothing about whether the encoder tracks cortex. What it does is remove one of the three readings of the

outrage null from the list of things this design cannot distinguish, which §6.2 takes up.

5.3.2 The confound, stated before the interpretation

Each category is drawn from exactly one source dataset: fear-activating from ISOT-fake, high outrage from SemEval-2019 Task 4, reward-hook from Webis-Clickbait-17, and neutral from PubMed and ISOT-true in equal parts. Category and provenance are therefore confounded by construction, and the separation reported above is a separation *between corpora*. It cannot be attributed to framing rather than to the writing conventions, vocabulary and editorial process of the source.

Length is not the confounding variable. Word counts are matched across categories at 167.4, 163.5, 163.7 and 162.2 with standard deviations near 11, which the corpus design achieved deliberately. Provenance is the variable that was not controlled.

One piece of internal evidence argues against a pure source artifact. High outrage is drawn from its own distinct source and does not separate from neutral on the index. An effect driven only by provenance would be expected to move every category away from the baseline, not three of four in an ordered way. The evidence is weaker than that sentence alone suggests: §5.3.1 shows high outrage moving on both components and only the difference staying flat, so a provenance effect is excluded only if it is assumed to act on one ROI union at a time. This is suggestive, not decisive.

The design that would settle it is a corpus in which framing varies within a single source, or in which one source appears in more than one category. Neither exists here, so the attribution remains open and is carried into the discussion as a limitation rather than resolved by argument.

5.4 RQ III: PHASE STRUCTURE AND THE COUPLING

The mean-field treatment of Chapter 3 gives the free energy $F(m) = am^2 + bm^4 - hm$ with $a = (1 - \beta J)/2$ and $b = 1/12$, coefficients derived from the




Figure 5.4: Predicted response for one fear-activating item at every one of the 20,484 cortical vertices, on the pial surface. Vertices below the map’s 70th percentile are left unpainted so the anatomy remains readable; the threshold is stated in the figure. This is the encoder’s own output and not a smoothed or interpolated field.

self-consistency condition rather than asserted. The phase structure is presented across a range of α and is not evaluated at a fitted value.

5.4.1 The coupling is unidentified

The calibration was run three times: on two partial corpora and on the completed 400 items. On the partial runs the confidence interval comfortably included zero. On the full corpus it does not, and the bootstrap over 10,000 replicates places no mass below zero.

That is not a measurement of the coupling, for two reasons stated together because either alone would be misleading. The fit’s held-out R^2 is **negative**, so it predicts less well than the sample mean on data it has not seen. And the estimate scales inversely with βJ , which no measurement in this study fixes: halving the assumed coupling strength roughly doubles the fitted value. A single number would therefore report a choice of βJ rather than a property of media.

The coupling is accordingly reported as unidentified, and no value is quoted.




Figure 5.5: Free energy at each category’s measured mean, swept over α . The coupling is swept and never fitted, and no value of $\hat{\alpha}$ is quoted.

5.4.2 The bound

What the corpus does supply is the observable’s measured spread, $\Delta X = 0.1241$ over the range -0.0698 to $+0.0543$. Chapter 3’s bound,

$$\alpha \geq \frac{h_c(\beta J)}{\Delta X}, \quad (5.3)$$

converts that spread into a statement about what the coupling would have to be:

βJ	$h_c(\beta J)$	α required
1.102	0.02101	0.1693
1.496	0.20541	1.6553
2.000	0.53284	4.2938

Table 5.7: The field bound at three values of the reduced coupling. The measured spread $\Delta X = 0.1241$ converts each critical field into the coupling any proposal of this mechanism would require.

This is a necessary condition and not a sufficient one. Clearing it would not establish that media of this measured spread moves populations; falling below it excludes the mechanism within the model’s assumptions. The value of the bound is that it requires no fit: it uses the measured spread and the model’s own spinodal, and it applies to any candidate observable proposed as a field.

5.5 STATISTICAL POWER

The design’s sensitivity was computed before the corpus was analysed and is reported alongside every result above, so that a reader can see what this study

could and could not have found. At $\alpha = 0.05$ and 80% power, with four categories of 100 and 300 against 100 for the classifier:

Smallest detectable effect	Value
Category separation, η^2	0.0268
Equivalently, Cohen's f	0.1659
Classifier AUC	0.5916
Equivalent separation d	0.3277

Table 5.8: Smallest effects this design could detect at $\alpha = 0.05$ and power 0.80, fixed before the corpus was analysed.

Both observed effects clear these thresholds, and power at the observed values is 1.0 for the separation and 0.961 for the AUC. Had either fallen below, the correct reading would have been that the design could not have detected an effect of that size, not that no effect exists.

5.6 RQ IV: OUT OF SCOPE

Research Question IV asked for per-modality mutual information across text, audio and video. The amendment reduced the corpus to 400 text items and dropped the audio and video arms, so the question as posed cannot be answered by this study. The narrowed text-only version was not reached within the project timeline. It is recorded here as out of scope rather than attempted and reported thinly.

5.7 SUMMARY OF THE ANSWERS

1. **RQ I.** The ratio form of the index is undefined for 17.25% of real content and is replaced by the signed difference. Against the pre-scan dataset-derived label the index achieves AUC 0.6274, above the detectable floor and in the predicted direction, but modest. A TF-IDF baseline on the same rows reaches 0.9758, so the index is not the better detector; a sentiment baseline does not clear chance.

2. **RQ II.** The four categories separate, $F = 15.779$, $p = 1.035 \times 10^{-9}$, $\eta^2 = 0.1068$. The effect is concentrated in fear-activating content ($d = +0.939$); high outrage does not separate on the index at all, but raises both component activations by almost the same amount ($d = +0.507$ and $d = +0.491$), so its flat difference is a symmetric response rather than an absence of one. Category is confounded with source dataset, so the separation cannot be attributed to framing.
3. **RQ III.** The coupling is unidentified and no value is quoted. The measured spread $\Delta X = 0.1241$ yields a lower bound on any coupling capable of driving a transition, reaching $\alpha \geq 4.2938$ at $\beta J = 2$.
4. **RQ IV.** Out of scope under the amendment.

Every category mean remains negative throughout. Whatever the categories differ in, none of them shows the affective-salience network leading in absolute terms.

Discussion

Chapter 5 reported that the four categories separate, that the separation is carried almost entirely by one of them, and that category is confounded with source dataset by construction. This chapter asks what can be concluded from that, and is organised around the order in which an examiner is entitled to doubt it.

6.1 WHAT THE SEPARATION CAN AND CANNOT SUPPORT

The measured effect, $\eta^2 = 0.1068$ at $p = 1.035 \times 10^{-9}$, is four times the smallest effect the design could detect. Taken on its own terms it is unambiguous: the index assigns systematically different values to the four groups of documents.

What it does not support is the sentence the proposal was written to test. Because each category is drawn from a single source dataset (§5.3.2), a separation between categories and a separation between corpora are the same measurement here. Fear-activating items are ISOT-fake items. Whatever distinguishes ISOT-fake from PubMed and ISOT-true — its framing, but equally its vocabulary, its house style, its editorial process, its era — is folded into the effect and cannot be separated from it by any analysis of this data.

This is a design limitation, not an analytical one. No amount of statistical work on 400 items assembled this way can attribute the difference to framing.

6.1.1 The evidence that argues against a pure source artifact

Three observations bear on it, none decisive.

First, high outrage does not separate on the index ($d = +0.030$, $p = 0.832$) despite coming from its own distinct source, SemEval-2019 Task 4. If the index were simply detecting provenance, every category drawn from a different corpus

should move away from the baseline. Three of four moving, in an ordered way, is a weaker pattern than pure source detection predicts. This argument has to be stated more carefully than it was before the components were split. Section 5.3.1 shows that high outrage does move away from the baseline on both A_{aff} and A_{del} taken separately; what it does not move is the difference. A pure provenance effect is therefore not excluded, but to produce the observed pattern it would have to shift both ROI unions by nearly identical amounts for one source while shifting only one union for another. That is a structured artifact rather than a generic one, and the observation is correspondingly weaker evidence than a flat outrage null would have been.

Second, length is controlled. Word counts sit at 167.4, 163.5, 163.7 and 162.2 with standard deviations near 11. The most obvious surface confound was designed out.

Third, the ordering is not arbitrary with respect to the hypothesis. Fear-activating content sits furthest from neutral and reward-hook content sits between. A source artifact carries no reason to order itself along the dimension the study proposed.

Against these stands the plain fact that the confound exists and is complete. The honest position is that the result is consistent with the hypothesis and equally consistent with a corpus artifact, and that this study cannot choose between them.

6.1.2 The experiment that would decide it

The confound is removable by design, not by analysis. Two routes are available to a successor:

1. **Within-source variation.** Draw all four categories from a single source, so provenance is held constant while framing varies. Feasible where a source contains both plainly-reported and emotively-framed material.
2. **Crossed design.** Ensure each source appears in more than one category, making source and category statistically separable so their contributions can be estimated independently.

Either would convert the present result from a suggestive to an interpretable one. Neither requires new instrumentation; both require rebuilding the corpus.

6.2 WHY HIGH OUTRAGE DOES NOT SEPARATE

The proposal expected outrage to separate most strongly, and it separated least. Three readings deserve stating, and unlike the confound of §5.3.2 the corpus can partly adjudicate between them.

The first is that the null is a property of the instrument rather than of the content. Outrage framing may engage deliberative and affective systems in more balanced measure than threat framing does, in which case a signed difference between the two networks is precisely the wrong instrument for detecting it: a large symmetric response and no response are indistinguishable in a difference.

The second is a property of the source. SemEval-2019 Task 4 labels hyperpartisan publishing at the article level. Hyperpartisan is not synonymous with outrage-framed, and a corpus selected on partisanship may contain a substantial fraction of items that are simply argumentative rather than emotionally charged.

The third is the encoder. TRIBE was trained on naturalistic film and podcast viewing, and is being asked to predict responses to synthesised speech from news text. There is no reason to expect its sensitivity to be uniform across framing styles that were absent from training.

The first reading is the one the data supports. Section 5.3.1 splits the index into the two activations it differences. High outrage raises A_{aff} against neutral at $d = +0.507$ and A_{del} at $d = +0.491$, both at $p < 10^{-3}$, and moves their difference by $d = +0.030$ at $p = 0.832$. That is the symmetric-response signature stated exactly: two effects of the same sign, of nearly the same size, cancelling in a difference. The comparison with fear-activating content makes the point sharper still, since fear moves A_{aff} at $d = +0.613$ and leaves A_{del} untouched ($d = -0.098$). The same instrument that registers a one-sided response registers nothing when the response is two-sided, which is the behaviour a difference is supposed to have.

This demotes the second and third readings rather than eliminating them. Neither predicts what was found. A corpus of merely argumentative items should

look like the neutral baseline on both components, and it does not: both are elevated, and A_{del} is elevated more for high outrage than for any other category. An encoder blind to the framing style should likewise leave both components near baseline. Each remains available as a partial account, since a source effect or an encoder effect could in principle raise both unions together, but neither is required to explain the null once the components are read separately, and the symmetric-response account explains it without further assumption.

What this does not do is rescue the design. The two components are computed from the same predicted activation over the same cortical proxy, so a common artifact would move them together and produce the same picture. The reading is supported, not established, and it is supported to the extent the encoder is trusted, which §6.3 takes up. It nonetheless changes what §5.3.1 permits to be said: the outrage null is a measured property of a difference between two responses that both moved, not an absence of response, and reporting it as an absence would have been wrong.

6.3 THE INSTRUMENT’S STANDING

Every value in this thesis is a prediction from a released encoder. No participant was scanned, no brain was measured, and the observable is a cortical proxy computed from predicted activation over two ROI unions.

That distinction is load-bearing here in a way it was not when the result looked like a null. A null in predicted activation is a weak statement whichever way the encoder behaves. A *detected separation* in predicted activation is only a statement about media if the encoder tracks cortex at all. The result reported in Chapter 5 therefore raises the stakes on a question this study cannot answer from its own data.

6.3.1 The checkpoint problem

The released checkpoint is loaded in its average-subject configuration, which `tribev2/demo_utils.py` forces at load time. A commercial laboratory has published an audit of that configuration reporting vertex-level $r = -0.0145$ against a

measured inter-subject ceiling of +0.0508: anti-correlated with real cortex, across four subjects and all seven Yeo networks [22, 21].

That audit is self-published by a company selling the per-subject alternative, uses four subjects and one stimulus, and reports magnitudes near 0.015 in either direction. It is a claim about sign, not about effect size, and it has not been independently replicated. It should not be accepted uncritically, and it cannot be dismissed either.

Work toward replicating it was begun within this project. Using the publicly released Algonauts 2025 responses [19], the leave-one-subject-out noise ceiling was measured directly from four subjects over all 1000 Schaefer parcels [20] of episode s01e01a, 592 timepoints:

$$r_{\text{ceiling}} = 0.1517, \quad 95\% \text{ CI } [0.1484, 0.1551]. \quad (6.1)$$

An earlier figure of 0.2247 is withdrawn. It was produced by a notebook cell that selected the first key in the response file, which is a different episode from the one under prediction, and whose orientation test could not fire on an array with fewer timepoints than parcels, so the correlation ran across the parcel axis and the reported count of 482 “parcels” was in fact the number of timepoints. The measurement now lives in a script that takes the episode as an argument and asserts orientation against the atlas size. This is the quantity any encoder must be judged against, and it is measured here rather than quoted. It is larger than the audit’s +0.0508 by construction and not by disagreement: averaging vertices within a parcel cancels independent noise and raises correlations, so a parcel-level ceiling and a vertex-level ceiling are not comparable numbers. The comparison against the checkpoint’s own predictions was not completed within this project’s timeline and is stated as unfinished rather than estimated.

6.3.2 What this means for building on foundation models

A technical report describing an architecture is not a warranty about the weights that are released under its name. The checkpoint used here is cortical-only, which

the proposal did not anticipate; it is loaded in a configuration the accompanying paper does not foreground; and its agreement with real cortex is contested in the literature. None of that is visible from the paper.

The practical lesson generalises beyond this project: anyone building a measurement on a released model should verify the checkpoint's properties from the artifact itself, as `scripts/verify_tribe_checkpoint.py` does here, rather than inheriting them from the publication.

6.4 THE COUPLING, AND WHAT THE BOUND DOES INSTEAD

The coupling between the index and the physics model's field is unidentified (§5.4): significant in sample, negative in held-out R^2 , and inversely dependent on a βJ that nothing in this study measures. Reporting a value would report a choice.

The bound is what survives. Because it requires only the observable's measured spread and the model's own spinodal, $\alpha \geq h_c(\beta J) / \Delta X$ converts $\Delta X = 0.1241$ into a statement no fit is needed to make: a coupling below 4.2938 at $\beta J = 2$ cannot push an all-to-all mean-field population across its spinodal, however the calibration is performed. This inverts the usual practice in the sociophysics literature, where the field is assigned and the population is shown to tip, without establishing that any real influence reaches the assigned magnitude.

That inversion is the contribution most likely to be reusable by others, because it applies to any proposed field observable and can be evaluated before a corpus is collected.

6.5 SESSION RELIABILITY, AND WHAT IT PERMITS

A second GPU session scanned all 400 items with identical text, identical code and identical ROI definitions. That pair is a test-retest measurement, and it is reported here because every effect in Chapter 5 otherwise rests on a single session with no statement of how much of an item's score is the item and how much is the session.

Agreement between the two sessions, over all 400 items:

	r	ICC(2,1)	$\text{sd}(\Delta)$	$\text{sd}(\Delta)/\text{sd}_{\text{between}}$
$\text{NAA}_{\text{signed}}$	0.8763	0.8725	0.01029	0.484
A_{aff}	0.8759	0.8713	0.01213	0.489
A_{del}	0.8576	0.8550	0.01271	0.521

Table 6.1: Agreement between the two scanning sessions over all 400 items. ICC is reported beside r because a session that shifted every score by a constant would still correlate perfectly.

The intraclass correlation is reported alongside r because a session that shifted every score by a constant would still correlate perfectly, and a shift is precisely what a second session can produce. The two agree closely here, so the disagreement is scatter rather than offset.

The separation replicates. Recomputing the four-category ANOVA on the second session over the same 400 items gives $\eta^2 = 0.0888$, $F = 12.861$, $p = 4.918 \times 10^{-8}$, against $\eta^2 = 0.1068$, $F = 15.779$, $p = 1.035 \times 10^{-9}$ for the first session over those items. Both clear the smallest effect this design could detect, $\eta^2 = 0.0268$. The separation reported in Chapter 5 is therefore not an artifact of one session. The second session’s effect is roughly a fifth smaller. Both sessions carry the same error, so the difference is within what sampling variation between two noisy runs produces.

Per-item direction is not reliable. The signed index is read as a direction, and 51 of 400 items, 12.8%, reverse that direction between sessions. A statement that a particular item is affective-leading is therefore not supportable by this instrument at this precision, and none is made: no per-item verdict appears in the analysis, in the figures or on the public interface. Group means over 100 items average this scatter down by an order of magnitude, which is why the category-level result survives while the item-level one does not.

The reversal rate is the reported scatter, not a separate defect. A reversal count means little until it is compared against what the measured noise already predicts, so the comparison is made here. Three mechanisms could move an item between sessions: the text is re-synthesised to speech each run, with a documented fallback to a different TTS model when the primary service is unreachable; the audio is force-aligned to obtain word events; and the encoder runs on a GPU.

Separating those three requires a controlled re-run and is not possible from the two session files. What is possible is asking whether they need separating at all.

The paired standard deviation gives a per-session error of $\sigma = \text{sd}(\Delta)/\sqrt{2} = 0.00728$, against a between-item standard deviation of 0.01913 once that error variance is removed from the variance of the session means. Under a plain additive-error model at that σ , with each item's session-free value estimated by shrinking its session mean toward the grand mean so that no item's own noise is used to predict its own reversal, the expected number of reversals over 400 items is 55.3, with a 95% interval of [44, 67] across 20,000 simulations. The observed 51 sits inside it, at $P(\text{simulated} \geq \text{observed}) = 0.79$ (`scripts/reversal_diagnosis.py`).

The reversals are also where that model says they should be. Items that reverse have a mean absolute index of 0.00320 against 0.02159 for items that do not ($p = 1.4 \times 10^{-25}$), and they are distributed evenly across the four categories ($\chi^2 = 2.764$, $p = 0.429$). The pattern is items whose value sits close enough to zero that the sign is undetermined at this precision, not a subset of content the instrument handles badly.

This changes what the controlled re-run is for. It is not a search for a defect, because the reversal count does not exceed what a single noise scale already explains; it is a measurement of how σ divides between synthesis, alignment and GPU nondeterminism, which decides whether σ can be reduced at all. It also settles the standing of per-item claims within this design rather than deferring them: with $\sigma = 0.00728$ against a between-item spread of 0.01913, a per-item direction is unavailable at any level of engineering care that leaves σ where it is. Reaching it needs either a lower-variance pipeline or repeated scans per item, and repeated scans are the cheaper of the two, since averaging k sessions divides σ by $\sqrt{k}$.

Both sessions are reported as measured. Neither is corrected toward the other and the two are not averaged, because an average of two sessions would carry a precision that neither run has.

Scan order does not carry the separation. The corpus took more than one session to scan, which raises the question of whether position in the scan, and

therefore session, is doing the work attributed to category. It is not, and the reason is that items were scanned interleaved rather than in category blocks: across thirds of the scan order the four categories are balanced ($\chi^2 = 3.139, p = 0.7913$). There is a real drift, with mean NAA_{signed} differing across thirds at $\eta^2 = 0.0167, p = 0.0353$. Because it is orthogonal to category by construction, it adds scatter rather than bias: removing the scan-third means before the category ANOVA moves the effect from $\eta^2 = 0.1068$ to $\eta^2 = 0.1048, p = 1.588 \times 10^{-9}$. Had the corpus been scanned one category at a time, this drift would have been indistinguishable from the finding.

6.6 LIMITATIONS

Stated here in full rather than distributed through the text.

1. **Category is confounded with source dataset.** The central limitation. Discussed in §5.3.2 and above.
2. **Single-session scores are noisy.** Test-retest ICC is 0.8725 and the session-to-session standard deviation is 48.4% of the between-item standard deviation (§6.5). Category-level conclusions survive this; per-item ones do not, and the 12.8% sign-reversal rate is what that noise scale predicts rather than an additional fault, so per-item direction is out of reach at this σ rather than pending a fix.
3. **Predicted, not measured, activation.** No brain was scanned. Whether the encoder tracks cortex is unresolved and contested.
4. **Cortical proxy, not a subcortical measurement.** The checkpoint is cortical-only. The structures the original proposal named cannot be measured with it, and no claim about them is made anywhere in this document.
5. **Average-subject configuration.** Forced by the loader. Whatever individual variation the model can express is unavailable.
6. **Out-of-distribution stimuli.** The encoder was trained on naturalistic audio-visual material and is fed synthesised speech from written text.

7. **The ratio form fails on 17.25% of items**, and the reformulation to a signed difference, while necessary, changes what the index means: it measures which network leads, not by what proportion.
8. **English-language corpus from five specific sources, none of them Kenyan or African.** A search for openly licensed Kenyan or African news data carrying usable labels returned nothing suitable, so every item in the corpus originates from American or European publishing. The measured spread is a property of that corpus, not of media in general, and the bound inherits the limitation. Whether the instrument behaves the same way on media a Kenyan reader actually encounters is untested, and assembling a labelled local corpus is the most obvious extension of this work.
9. **Mean-field physics.** All-to-all coupling, one homogeneous population, a static field, binary opinions. Network structure and heterogeneity generally lower h_c , so on a real network the bound can overstate the coupling required. It is a statement about the all-to-all model and not a universal one.
10. **The interpretation is a reverse inference.** Naming the two ROI unions affective-salience and deliberative-control, and reading their difference as a statement about which mode of processing a text engages, runs the inference backwards: from a spatial pattern to a cognitive process. Regions do not map one-to-one onto processes. Amygdala, insula and anterior cingulate activate across threat, reward, interoception, salience of any valence and plain physiological arousal; prefrontal and parietal control regions activate across working memory, task switching, conflict and effort of every kind. The same difference between the two unions is therefore consistent with a great many things other than a reader being manipulated, and this study has no independent behavioural or self-report measure to break the tie. The limitation compounds with the third and fourth entries above: the pattern being interpreted is predicted rather than measured, over a cortical proxy for structures the naming invokes. What the index can be said to measure without reverse inference is narrow and is what Chapter 5 claims: a signed

difference between two predicted ROI-union activations, which separates four content categories and correlates with a pre-existing label.

11. **A lexical baseline separates the label better.** Section 5.2.3 reports TF-IDF with logistic regression at AUC 0.9758 against the index's 0.6274 on the same rows and the same label. Nothing in this thesis argues that the index is a better detector of manipulative content than counting words, and any use of it as a screening tool would have to justify itself on something other than classification accuracy.

6.7 WHAT WOULD BE REQUIRED NEXT

In order of what each would buy.

A crossed or within-source corpus would make the separation interpretable. It is the cheapest of these and the one that most changes what can be claimed.

Completing the validation would establish whether the instrument measures anything about brains. The noise ceiling is measured and the evaluation code is written and tested; what remains is running the checkpoint over a held-out stimulus and correlating. It needs no large compute.

A variance decomposition of the scanning pipeline would say whether σ can be lowered. One item re-scanned ten times with the text fixed and the synthesised audio held constant, then ten more with synthesis free, separates the contribution of speech synthesis from that of alignment and GPU nondeterminism. Section 6.5 establishes that this is a question about the size of σ rather than a search for a defect, since the reversal rate is already accounted for. It is the cheapest experiment on this list, needing twenty scans of one item.

A subcortical checkpoint would let the original proposal's question be asked as posed. That means the training grid, the four source fMRI corpora and compute well outside the scope of an undergraduate project, and it is recorded as future work rather than a plan.

6.8 SCREENING MEDIA DESIGNED TO HOLD ATTENTION

The proposal defence raised a specific use for an instrument of this kind: scan a piece of media and say whether it is built to stimulate rather than to inform, with children’s programming named as the case where it would matter most. The suggestion is worth answering directly, because the instrument built here is a partial answer and the gap between it and that application is concrete rather than vague.

What already carries over. The pipeline accepts video natively. Its video entry point, `InferenceService.predict_video`, takes a file, extracts the encoder’s own event representation and returns predictions per repetition time. That matters because the encoder was trained on naturalistic audiovisual material, so a children’s programme is closer to its training distribution than the synthesised speech from news text used throughout this thesis, which is an out-of-distribution stimulus (§6.3). The instrument would be on firmer ground there, not weaker.

The nearest evidence this corpus offers is the reward-hook category, which is clickbait, the closest available proxy for content engineered to capture attention. It separates from neutral at $d = +0.319$, $p = 0.025$: real, modest, and subject to the same source confound as everything else here.

What is missing, and it is not compute. Three things stand between this instrument and that application.

1. **The time axis is discarded.** `predict_video` returns an ($\text{TRs} \times \text{vertices}$) array which is then mean-pooled to a single map, and every result in this thesis is computed from that pooled map. Content designed to hold attention is a claim about *structure over time*: pacing, the timing of reward events, how quickly a lull is resolved. A single averaged value cannot express it, and averaging is exactly the operation that destroys it. The data is already produced and thrown away, so this is an analysis change rather than a new measurement.
2. **Behavioural ground truth.** “Designed to be addictive” is a statement about what content does to viewing behaviour, and this instrument predicts cortical

response, not behaviour. Establishing the link needs retention or watch-time data alongside the media, so the predicted response can be validated against something outside the model. Without it the instrument produces a number with no established relationship to the property being claimed.

3. **Per-item reliability.** §6.5 showed 12.8% of items reverse direction between sessions. Judging a single programme is precisely a per-item claim, and this instrument cannot support one at present. Group comparisons, for instance between platform-optimised and public-service children's programming, are within reach; a verdict on one show is not.

A study that would work. Two matched sets of children's programming differing on a single axis, drawn where possible from the same producers so provenance does not stand in for the property of interest; video passed through the encoder directly; the observable computed over time rather than pooled; and validation against retention data. That design answers the question the defence asked, and each of its four requirements is a direct response to a limitation established in this chapter.

Two cautions belong with it. Asserting that a specific programme is engineered to be addictive is a serious accusation about identifiable producers, and it needs an evidentiary bar far above what a predicted cortical asymmetry currently supplies. And the instrument would need the validation of §6.3 settled first: an unvalidated encoder applied to children's media would produce confident numbers about a sensitive subject with nothing underwriting them.

Conclusion

7.1 WHAT WAS PROMISED

The proposal committed to an instrument. It said so three times: that the primary deliverable is a functional Asynchronous Measurement Apparatus, that this is “not a proposed future application and not a conceptual description”, and that the project “proposes a thermometer, and delivers it as working, open-source code that can be run today”.

That framing governs how this work should be judged. A thermometer that reads a value nobody expected is still a delivered thermometer. Failing to ship it would have been the failure.

7.2 WHAT WAS DELIVERED

7.2.1 The apparatus

The instrument exists and runs. Given a passage of text it synthesises speech, obtains word timings, computes text and audio embeddings, passes them to the released encoder, and returns a predicted response over 20,484 cortical surface points, from which the index is computed. It is open source, it is deployed, and it was run end to end over 400 items.

Its outputs are reproducible: the corpus file, the analysis scripts, the power calculation and the physics all regenerate from the repository, and every number in Chapters 5 and 6 is produced by a script rather than transcribed.

7.2.2 The framework

The mean-field treatment of Chapter 3 derives its Landau coefficients from the self-consistency condition rather than asserting them, recovers the exact critical exponents numerically as a check on the solver, and locates the spinodal field $h_c(\beta J)$ at which bistability collapses.

From that follows the bound: for any observable X proposed as a field through $h = \alpha X$, no media-driven transition across the spinodal of the all-to-all mean-field model is possible unless

$$\alpha \geq \frac{h_c(\beta J)}{\Delta X}. \quad (7.1)$$

This inverts the usual procedure in the sociophysics literature, where the field is assigned by hand and the population is then shown to tip. It requires no fit, applies to any candidate observable, and can be evaluated before a corpus is collected.

7.2.3 The measurement

All 400 items were scanned and verified. The four categories separate at $\eta^2 = 0.1068$, $p = 1.035 \times 10^{-9}$, against a design powered to detect 0.0268. Against the labels the items carried before scanning, the index achieves AUC 0.6274, above the detectable floor of 0.5916 and in the predicted direction.

The separation is carried almost entirely by fear-activating content ($d = +0.939$). High outrage, the category the proposal expected to separate most, does not separate at all ($d = +0.030$, $p = 0.832$). Split into the two activations the index differences, high outrage raises both against neutral by nearly the same amount ($d = +0.507$ and $d = +0.491$, both $p < 10^{-3}$), so what looks like an absent response is a symmetric one that a difference cannot see.

A TF-IDF baseline on the same rows and the same label reaches AUC 0.9758. The index is not the better classifier and is not offered as one; it is offered as the field observable Chapter 3 requires.

Every category mean is negative. The deliberative-control network is predicted to respond more strongly than the affective-salience network for all four categories; they differ in how far below zero they sit, not in which network leads.

7.3 WHAT IS CLAIMED, AND WHAT IS NOT

Claimed. That the instrument works and is reproducible. That the index assigns systematically different values to these four groups of documents, at an effect size the design was powered to detect, with the power statement reported alongside. That the measured spread $\Delta X = 0.1241$ places a lower bound on any coupling capable of driving a transition, reaching $\alpha \geq 4.2938$ at $\beta J = 2$.

Not claimed. That the index detects manipulation, or detects it better than counting words does: it does not. That the separation is due to framing rather than to source dataset: category and provenance are confounded by construction and this design cannot distinguish them. That the coupling has been measured: it is unidentified, and no value of $\hat{\alpha}$ appears anywhere in this document. That anything here measures a brain: every value is a prediction from a released encoder, the observable is a cortical proxy, and no participant was scanned.

The gap between those two lists is the honest content of this work.

7.4 CONTRIBUTION

Three things, in the order they are likely to be useful to someone else.

The bound. A necessary condition on the coupling between any content observable and an opinion-dynamics field, requiring only the observable’s measured spread and the model’s own spinodal. It is the part of this work that transfers beyond the instrument that produced it.

The instrument, with its limitation documented. A working, open pipeline from text to predicted cortical response, together with the finding that the released checkpoint is cortical-only, that the loader forces its average-subject configuration, and that its agreement with real cortex is contested. A technical report describing an architecture is not a warranty about released weights, and this project’s experience is a concrete instance of that.

The measurement and its confound. A separation reported with the design’s sensitivity attached and its central weakness stated before its interpretation, together with the specific experiment that would resolve it.

7.5 FUTURE WORK

Remove the confound. A corpus in which framing varies within a single source, or in which sources are crossed with categories, would make the separation interpretable. It needs no new instrumentation, only a rebuilt corpus, and it is the single change that most alters what can be claimed.

Finish the validation. Whether the encoder tracks cortex at all is unresolved and determines what the measurement means. The noise ceiling has been measured directly from public data at 0.1517, CI [0.1484, 0.1551], and the evaluation code is written and tested. What remains is running the checkpoint over a held-out stimulus and correlating, which needs no large compute.

A subcortical checkpoint. The structures the original proposal named cannot be measured with a cortical-only model. Training one means the full grid, the source fMRI corpora, and compute outside the scope of an undergraduate project.

7.6 CODE, DATA AND THE WORKING INSTRUMENT

The pipeline, the corpus builder, the analysis scripts and the figure scripts are public at <https://github.com/brn-mwai/monarch>. Every number in this document is written by a named script into a JSON or CSV artifact and read from there rather than transcribed, so any figure or table can be regenerated from the repository alone.

The instrument itself runs at <https://monarch-4iy.pages.dev>. It carries the complete 400-item corpus with each item's text, its pre-scan labels and its measured values, and it draws the predicted response on the cortical surface for the items whose per-vertex maps were kept. It answers questions about the corpus from a generated fact sheet and declines anything the data does not cover, including any request for a coupling value.

7.7 CLOSING

This work set out to make an invisible property of media measurable, and to say honestly what the measurement supports. It delivers a working instrument, a

bound that constrains any future proposal of this kind, and a result that is real, modest, and confounded in a way this document states plainly rather than leaves to be discovered.

The physics is the part that survives whatever becomes of the encoder. Whether or not the released checkpoint predicts any brain, the requirement that $\alpha \geq h_c(\beta J)/\Delta X$ holds, within the all-to-all mean-field model, for every observable anyone proposes as a field, and the burden it places on such proposals does not depend on this instrument being the right one.

The longer aim behind the project, set out in §1.6, is a tool that a reader rather than a researcher could use to ask whether something they are about to watch is built to manipulate them or to hold their attention. This document is the first phase of that and makes no claim to be more. It leaves an instrument, a characterised measurement, an honest account of where that measurement fails, and a bound that does not depend on any of the above. What it does not leave is a tool anyone should use to judge a particular article or programme, and the four reasons why are measured in these chapters rather than guessed at. Stating the destination is not the same as arriving at it, and the difference between them is most of what this thesis reports.

Bibliography

- [1] C. Castellano, S. Fortunato, V. Loreto, Statistical physics of social dynamics, *Rev. Mod. Phys.* 81 (2009) 591–646. doi:10.1103/RevModPhys.81.591.
- [2] S. Galam, Sociophysics: A review of Galam models, *Int. J. Mod. Phys. C* 19 (2008) 409–440. doi:10.1142/S0129183108012297.
- [3] K. Sznajd-Weron, J. Sznajd, Opinion evolution in closed community, *Int. J. Mod. Phys. C* 11 (2000) 1157–1165. doi:10.1142/S0129183100000936.
- [4] D. Stauffer, Social applications of two-dimensional Ising models, *Am. J. Phys.* 76 (2008) 470–473. doi:10.1119/1.2779882.
- [5] P. Mullick, P. Sen, Sociophysics models inspired by the Ising model, *Eur. Phys. J. B* 98 (2025) 206. arXiv:2506.23837. doi:10.1140/epjb/s10051-025-01053-7.
- [6] N. Crokidakis, Effects of mass media on opinion spreading in the Sznajd sociophysics model, *Physica A* 391 (2012) 1729. arXiv:1111.5750. doi:10.1016/j.physa.2011.11.038.
- [7] Azhari, R. Muslim, The external field effect on opinion formation based on the majority rule and the q -voter models on the complete graph, *Int. J. Mod. Phys. C* (2023) 2350088. arXiv:2305.06051. doi:10.1142/S0129183123500882.
- [8] F. Freitas, A. R. Vieira, C. Anteneodo, Imperfect bifurcations in opinion dynamics under external fields, *J. Stat. Mech.* (2020) 024002. doi:10.1088/1742-5468/ab6848.
- [9] A. Sîrbu, V. Loreto, V. D. P. Servedio, F. Tria, Opinion dynamics: models, extensions and external effects, in: *Participatory Sensing, Opinions and Collective Awareness*, Springer, 2016, pp. 363–401. arXiv:1605.06326. doi:10.1007/978-3-319-25658-0_17.

- [10] J. Korbel, R. Dahdoul, S. Thurner, Empirical validation of the polarization transition in a double-random field model of elections, *Phys. Rev. Lett.* 136 (2026) 127402. arXiv:2510.00612. doi:10.1103/9gjj-1df6.
- [11] S. d’Ascoli, J. Rapin, Y. Benchetrit, H. Banville, J.-R. King, TRIBE: TRImodal Brain Encoder for whole-brain fMRI response prediction, arXiv:2507.22229 (2025).
- [12] S. d’Ascoli, J. Rapin, Y. Benchetrit, T. Brooks, K. Begany, J. Raugel, H. Banville, J.-R. King, A foundation model of vision, audition, and language for in-silico neuroscience, arXiv:2605.04326 (2026).
- [13] N. Goldenfeld, *Lectures on Phase Transitions and the Renormalization Group*, Addison-Wesley, 1992; reprinted CRC Press, 2018. doi:10.1201/9780429493492.
- [14] H. Ahmed, I. Traore, S. Saad, Detection of online fake news using n-gram analysis and machine learning techniques, in: *Intelligent, Secure, and Dependable Systems in Distributed and Cloud Environments (ISDDC 2017)*, *Lecture Notes in Computer Science* 10618, Springer, 2017, pp. 127–138. doi:10.1007/978-3-319-69155-8_9.
- [15] J. Kiesel, M. Mestre, R. Shukla, E. Vincent, P. Adineh, D. Corney, B. Stein, M. Potthast, SemEval-2019 Task 4: Hyperpartisan news detection, in: *Proceedings of the 13th International Workshop on Semantic Evaluation*, 2019, pp. 829–839. doi:10.18653/v1/S19-2145.
- [16] M. Potthast, T. Gollub, K. Komlossy, S. Schuster, M. Wiegmann, E. P. Garces Fernandez, M. Hagen, B. Stein, Crowdsourcing a large corpus of clickbait on Twitter, in: *Proceedings of the 27th International Conference on Computational Linguistics (COLING)*, 2018, pp. 1498–1507. <https://aclanthology.org/C18-1127/>.
- [17] National Center for Biotechnology Information, PubMed, accessed through the Entrez Programming Utilities, <https://pubmed.ncbi.nlm.nih.gov>.

- [18] C. J. Hutto, E. Gilbert, VADER: A parsimonious rule-based model for sentiment analysis of social media text, in: Proceedings of the Eighth International AAAI Conference on Weblogs and Social Media (ICWSM), 2014, pp. 216–225. doi:10.1609/icwsm.v8i1.14550.
- [19] A. T. Gifford, D. Bersch, M. St-Laurent, et al., The Algonauts Project 2025 Challenge: How the human brain makes sense of multimodal movies, arXiv:2501.00504 (2025).
- [20] A. Schaefer, R. Kong, E. M. Gordon, T. O. Laumann, X.-N. Zuo, A. J. Holmes, S. B. Eickhoff, B. T. T. Yeo, Local-global parcellation of the human cerebral cortex from intrinsic functional connectivity MRI, *Cereb. Cortex* 28 (2018) 3095–3114. doi:10.1093/cercor/bhx179.
- [21] B. T. T. Yeo et al., The organization of the human cerebral cortex estimated by intrinsic functional connectivity, *J. Neurophysiol.* 106 (2011) 1125–1165. doi:10.1152/jn.00338.2011.
- [22] The Sapient Company, The Faceless Brain, self-published technical report, 2026, <https://thesapientcompany.com/research>.